ISIT312 Big Data Management

Hive Programming

Dr Fenghui Ren

School of Computing and Information Technology -
University of Wollongong

Hive Programming

Outline

Data Selection and Scope

To query data **Hive** provides `SELECT` statement

Typically `SELECT` statement projects the rows satisfying the query conditions specified in the `WHERE` clause and returns the result set

`SELECT` statement is usually used with `FROM`, `DISTINCT`, `WHERE`, and `LIMIT` keywords

```
                                        SELECT statement with LIMIT clause
SELECT C_NAME, C_PHONE
FROM customer
WHERE C_ACCTBAL > 0
LIMIT 2;
```

Data Selection and Scope

Multiple `SELECT` statements can be combined into complex queries using nested queries or subqueries

Subqueries can use Common Table Expressions (CTE) in the format of `WITH` clause

When using subqueries, an alias should be given for the subquery

```
                                                    WITH clause
WITH cord AS ( SELECT *
               FROM customer JOIN orders
                             ON c_custkey = o_custkey)
SELECT c_name, c_phone, o_orderkey, o_orderstatus
FROM cord;
```

Data Selection and Scope

Multiple `SELECT` statements can be combined into complex queries using nested queries or subqueries

Nested queries can use `SELECT` statement wherever a table is expected or a scalar value is expected

```
                                                        Nested query
SELECT c_name, c_phone, o_orderkey, o_orderstatus
FROM ( SELECT *
       FROM customer JOIN orders
                ON c_custkey = o_custkey) cord
```

Data Selection and Scope

When inner join is performed between multiple tables the MapReduce jobs are created to process data in HDFS

It is recommended to put the big table right at the end for better because the last table in the sequence is streamed through the reducers where the others are buffered in the reducer by default

```
                                                              Inner join
SELECT /*+ STREAMTABLE(lineitem) */ c_name, o_orderkey, l_linenumber
FROM customer JOIN orders
              ON c_custkey = o_custkey
              JOIN lineitem
              ON l_orderkey = o_orderkey;
```

Data Selection and Scope

Outer join (left, right, and full) and cross join preserve their HQL semantics

Map join means that join is computed only by map job without reduce job

In map join all data are read from a small table to memory and broadcasted to all maps

During map phase each row in from a big table is compared with the rows in small tables against join conditions

Join performance is improved because there is no reduce phase

Data Selection and Scope

```
SELECT /*+ MAPJOIN(orders) */ c_name, c_phone, o_orderkey, o_orderstatus
FROM customer JOIN orders
            ON c_custkey = o_custkey;
```

Hive automatically converts the `JOIN` to `MAPJOIN` at runtime when `hive.auto.convert.join` setting is set to `true`

Bucket map join is a special type of `MAPJOIN` that uses bucket columns in join condition.

Then instead of fetching the whole table bucket map join only fetches the required bucket data.

A variable `hive.optimize.bucketmapjoin` must be set to true to enable bucket map join

Data Selection and Scope

Hive supports `LEFT SEMI JOIN`

```
                                                Left semi join
SELECT c_name, c_phone
FROM customer LEFT SEMI JOIN orders
              ON c_custkey = o_custkey;
```

In `LEFT SEMI JOIN` the right-hand side table should only be referenced in the join condition and not in `WHERE` or `SELECT` clauses

Data Selection and Scope

Hive supports `UNION ALL` it does not support `INTERSECT` and `MINUS` operations

```
SELECT p_name                                    UNION ALL operation
FROM PART
UNION ALL
SELECT c_name
FROM CUSTOMER;
```

`INTERSECT` operation can be implemented as `JOIN` operation

`MINUS` operation can be implemented as `LEFT OUTER JOIN` operation with `IS NULL` condition in `WHERE` clause

Hive Programming

Outline

Data Manipulation

LOAD statement can be used to load data to Hive tables from local file system or from HDFS

Load data to Hive table from a local file

```
                                          Loading data from a local file
LOAD DATA LOCAL INPATH '/local/home/janusz/HIVE-EXAMPLES/TPCHR/part.txt'
OVERWRITE INTO TABLE part;
```

Load data to Hive partitioned table from a local file

```
                              Loading data into partitioned table from a local file
LOAD DATA LOCAL INPATH '/local/home/janusz/HIVE-EXAMPLES/TPCHR/part.txt'
OVERWRITE INTO TABLE part PARTITION
(P_BRAND='GoldenBolts');
```

LOCAL keyword determines a location of the input files

Data Manipulation

Load HDFS data to the Hive table using the default system path

```
                                                    Loading dta from HDFS
LOAD DATA INPATH '/user/janusz/part.txt'
OVERWRITE INTO TABLE part;
```

Load HDFS data to the Hive table using using full URI

```
                                                    Loading data from HDFS
LOAD DATA INPATH 'hdfs://10.9.28.14:8020/user/janusz/part.txt'
OVERWRITE INTO TABLE part;
```

If LOCAL keyword is not specified, the files are loaded from the full URI specified after INPATH or the value from the fs.default

OVERWRITE keyword decides whether to append or replace the existing data in the target table/partition

Data Manipulation

`EXPORT` and `IMPORT` statements are available to support the import and export of data in HDFS for data migration or backup/restore purposes

`EXPORT` statement exports both data and metadata from a table or partition

```
EXPORT TABLE part TO '/user/tpchr/part'
```

Metadata is exported to a file called `_metadata`

```
-rwxr-xr-x 3 janusz supergroup 2739 2017-07-09
14:37 /user/tpchr/part/_metadata
drwxr-xr-x - janusz supergroup 0 2017-07-09
14:37 /user/tpchr/part/p_brand=GoldenBolts
```

After `EXPORT` the exported files can be copied to other Hive instances or to other HDFS clusters

Data Manipulation

IMPORT statement imports files exported from other HIVE instances into an internal table

```
IMPORT table new_part FROM '/user/tpchr/part';
```
HQL

An imported table is located in a default HIVE location in HDFS

```
drwxrwxr-x - janusz supergroup 0 2017-07-09
14:56 /user/hive/warehouse/new_part
```
Importing data from HDFS

IMPORT EXTERNAL statement imports a file exported from other HIVE instances into an external table

Importing external table from HDFS

```
IMPORT EXTERNAL table new_extpart FROM '/user/tpchr/
part';
```

An imported table is located in a default HIVE location in HDFS

Contents of HDFS

```
drwxrwxr-x - janusz supergroup 0 2017-07-09
15:04 /user/hive/warehouse/new_extpart
```

Data Manipulation

ORDER BY sorts the results of SELECT statement

An order is maintained across all of the output from every reducer and global sort is performed using only one reducer

```
SELECT p_partkey, p_name                          ORDER BY clause
FROM part
ORDER BY p_name ASC;
```

SORT BY does the same job as ORDER BY and indicates which columns to sort when ordering the reducer input records

SORT BY completes sorting before sending data to the reducer

SORT BY statement does not perform a global sort and only makes sure data is locally sorted in each reducer

```
SET mapred.reduce.tasks = 2;                       SORT BY clause
SELECT p_partkey, p_name
FROM part
SORT BY p_name ASC;
```

Data Manipulation

When `DISTRIBUTE BY` clause is applied rows with matching column values are partitioned by the same reducer

```
                                              DISTRIBUTE BY clause
SELECT p_partkey, p_name FROM part
DISTRIBUTE BY p_partkey
SORT BY p_name;
```

`DISTRIBUTE BY` clause is similar to `GROUP BY` in relational systems in terms of deciding which reducer is used to distribute the mapper

When using with `SORT BY`, `DISTRIBUTE BY` must be specified before the `SORT BY` statement

Data Manipulation

CLUSTER BY clause is a shorthand operator to perform

DISTRIBUTE BY and SORT BY operations on the same group of columns.

```
                                              CLUSTER BY clause
SELECT p_partkey, p_name
FROM part
CLUSTER BY p_name;
```

ORDER BY performs a global sort, while CLUSTER BY sorts in each distributed group

To fully utilize all the available reducers we can do CLUSTER BY first and then ORDER BY

```
                                              CLUSTER BY clause
SELECT p_partkey, p_name
FROM part
CLUSTER BY p_name
ORDER BY p_name;
```

Hive Programming
Outline

Data Aggregation and Sampling

Hive supports several aggregation functions, analytic functions working with `GROUP BY` and `PARTITION BY`, and windowing clauses

Hive supports advanced aggregation by using `GROUPING SETS,` `ROLLUP, CUBE`, analytic functions, and windowing

Basic aggregation uses `GROUP BY` clause and aggregation functions

```
                                              GROUP BY clause
SELECT p_type, count(*)
FROM part
GROUP BY p_type;
```

To aggregate into sets a function `collect_set` can be used

```
                              GROUP BY clause with collect_set function
SELECT p_type, collect_set(p_name), count(*)
FROM part
GROUP BY p_type;
```

Data Aggregation and Sampling

GROUPING SETS clause implements advanced multiple GROUP BY operations against the same set of data

```
SELECT p_type, p_name, count(*)
FROM part
GROUP BY p_type, p_name
GROUPING SETS ( (p_type), (p_name) );
```

ROLLUP clause allows to calculate multiple levels of aggregations across a specified group of dimensions

```
SELECT p_type, p_name, count(*)
FROM part
GROUP BY p_type, p_name WITH ROLLUP;
```

CUBE clause allows to create aggregations over all possible subsets of attributes in a given set

```
SELECT p_type, p_name, count(*)
FROM part
GROUP BY p_type, p_name WITH CUBE;
```

Data Aggregation and Sampling

`GROUPING__ID` function works as an extension to distinguish entire rows from each other

```
GROUPING__ID function

SELECT GROUPING__ID, p_type, p_name, count(*)
FROM part
GROUP BY p_type, p_name WITH CUBE
ORDER BY grouping__id;
```

`HAVING` can be used for the conditional filtering of `GROUP BY` results

```
GROUPING__ID function

SELECT GROUPING__ID, p_type, p_name, count(*)
FROM part
GROUP BY p_type, p_name WITH CUBE
HAVING count(*) > 1
ORDER BY grouping__id;
```

Data Aggregation and Sampling

Analytic functions scan multiple input rows to compute each output value

Analytic functions are usually used with `OVER, PARTITION BY, ORDER BY`, and the windowing specification

Analytic functions operate on windows where the input rows are ordered and grouped using flexible conditions expressed through an `OVER PARTITION` clause

Syntax is the following

```
                                    Syntax of analytic functions
function (arg1,..., argn)
OVER ([PARTITION BY <...>]
[ORDER BY <....>] [])
```

For standard aggregation `function (arg1,..., argn)` can be either `COUNT(), SUM(), MIN(), MAX(),` or `AVG()`

Data Aggregation and Sampling

Typical aggregations implemented as analytic functions in the following way

```
                                              PARTITION BY clause
SELECT p_name,
       COUNT(*) OVER (PARTITION BY p_name)
FROM PART;
```

Other analytic functions are used as follows

```
                                              ORDER BY clause
SELECT l_orderkey, l_partkey, l_quantity,
       RANK() OVER (ORDER BY l_quantity),
       DENSE_RANK() OVER (ORDER BY l_quantity)
FROM lineitem;
```

```
                                              PARTITION BY clause
SELECT l_orderkey, l_partkey, l_quantity,
       RANK() OVER (PARTITION BY l_orderkey ORDER BY l_quantity),
       DENSE_RANK() OVER (PARTITION BY l_orderkey ORDER BY l_quantity)
FROM lineitem;
```

Data Aggregation and Sampling

More analytic functions ...

```
                                                    PARTITION BY clause
SELECT l_orderkey, l_partkey, l_quantity,
       FIRST_VALUE(l_quantity) OVER (PARTITION BY l_orderkey ORDER BY l_quantity),
       LAST_VALUE(l_quantity) OVER (PARTITION BY l_orderkey ORDER BY l_quantity)
FROM lineitem;;
```

```
                                                    PARTITION BY clause
SELECT l_orderkey, l_partkey, l_quantity,
       MAX(l_quantity) OVER (PARTITION BY l_orderkey ORDER BY l_partkey
                            ROWS BETWEEN 2 PRECEDING AND CURRENT ROW)
FROM lineitem;
```

```
                                                    PARTITION BY clause
SELECT l_orderkey, l_partkey, l_quantity,
       MAX(l_quantity) OVER (PARTITION BY l_orderkey ORDER BY l_partkey
                            ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING)
FROM lineitem;
```

Data Aggregation and Sampling

When data volume is extra large we can use a subset of data to speed up data analysis.

Random sampling uses the `RAND()` function and `LIMIT` clause to get the samples of data

```
                                              DISTRIBUTE BY clause

SELECT *
FROM lineitem DISTRIBUTE BY RAND() SORT BY RAND() LIMIT 5;
```

`DISTRIBUTE` and `SORT` clauses are used here to make sure the data is also randomly and efficiently distributed among mappers and reducers

Bucket table sampling is a special sampling optimized for bucket tables

```
                                                  Bucket sampling

SELECT *
FROM customer TABLESAMPLE(BUCKET 3 OUT OF 8 ON rand());
```

Data Aggregation and Sampling

Block sampling allows to randomly pick up N rows of data, percentage (n percentage) of data size, or N byte size of data

```
                                                    Block sampling
SELECT *
FROM lineitem TABLESAMPLE(4 ROWS);
```

```
                                                    Block sampling
SELECT *
FROM lineitem TABLESAMPLE(50 PERCENT);
```

```
                                                    Block sampling
SELECT *
FROM lineitem TABLESAMPLE(20B);
```

References

Gross C., GuptaA., Shaw S., Vermeulen A. F., Kjerrumgaar D., Practical Hive: A guide to Hadoop's Data Warehouse System, Apress 2016, Chapter 4 (Available through UOW library)

Lee D., Instant Apache Hive essentials how-to: leverage your knowledge of SQL to easily write distributed data processing applications on Hadoop using Apache Hive, Packt Publishing Ltd. 2013 (Available through UOW library)

Apache Hive TM, https://hive.apache.org/