

ISIT312 Big Data Management

HBase Data Model

Dr Fenghui Ren

School of Computing and Information Technology -
University of Wollongong

HBase Data Model

Outline

[Background](#)

[Logical view of data](#)

[Design fundamentals](#)

[Physical implementation](#)

Background

Hbase is open source distributed database based on a data model of Google's BigTable

HBase provides a BigTable view of data stored in HDFS

HBase is also called as Hadoop DataBase

HBase still provides a tabular view of data however it is also very different from the traditional relational data model

HBase data model is a sparse, distributed, persistent multidimensional sorted map

It is indexed by a row key, column key, and timestamp

HBase Data Model

Outline

[Background](#)

[Logical view of data](#)

[Design fundamentals](#)

[Physical implementation](#)

Logical view of data

HBase organizes data into **tables**

HBase **table** consists of **rows**

Each **row** is uniquely identified by a **row key**

Data within a **row** is grouped by a **column family**

Column families have an important impact on the **physical implementation** of **HBase table**

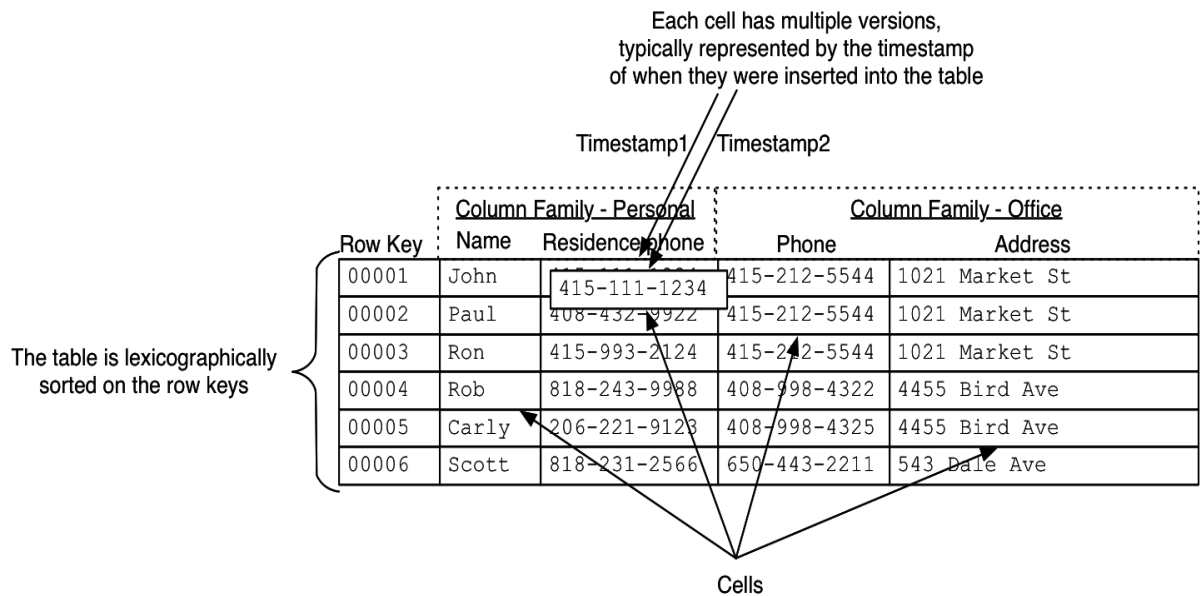
Every **row** has the same **column families** although some of them can be **empty**

Data within a **column family** is addressed via its **column qualifier**, or simply, **column name**

Hence, a combination of **row key**, **column family**, and **column qualifier** uniquely identifies a **cell**

Values in cells do not have a data type and are always treated as **sequences of bytes**

Logical view of data



Values within a cell have multiple versions

Versions are identified by their version number, which by default is a timestamp when the cell was written

Logical view of data

If a **timestamp** is not determined at write time, then the **current timestamp** is used

If a **timestamp** is not determined during a read, the **latest one** is returned

The **maximum allowed number of cell value versions** is determined for each **column family**

The **default number of cell versions** is three

Logical view of data

A view of **HBase table** as a nested structure

```
{ "Row-0001":  
  { "Home":  
    { "Name":  
      { "timestamp-1": "James" }  
    }  
    "Phones":  
      { "timestamp-1": "2 42 214339"  
        "timestamp-2": "2 42 213456"  
        "timestamp-3": "+61 2 4567890" }  
    }  
    "Office":  
      { "Phone":  
        { "timestamp-4": "+64 345678" }  
      }  
      "Address":  
        { "timestamp-5": "10 Ellenborough Pl" }  
      }  
    }  
  }  
}
```

HBase Table

Logical view of data

A view of **HBase table** as a nested structure

```
{ "Row-0002":  
  { "Home":  
    { "Name":  
      { "timestamp-6": "Harry" }  
    }  
    "Phones":  
      { "timestamp-7": "2 42 214234" }  
    }  
  "Office":  
    { "Phone":  
      { "timestamp-8": "+64 345678" }  
    }  
    "Address":  
      { "timestamp-9": "10 Bong Bong Rd"  
        "timestamp-10": "23 Victoria Rd" }  
    }  
  }  
}
```

HBase Table

Logical view of data

A **key** can be **row key** or a combination of a **row key**, **column family**, **qualifier**, and **timestamp** depending on what supposed to be retrieved

If all the **cells** in a row are of interest then a **key** is a **row key**

If only specific **cells** are of interest, the appropriate **column families** and **qualifiers** are a part of a **key**

HBase Data Model

Outline

[Background](#)

[Logical view of data](#)

[Design fundamentals](#)

[Physical implementation](#)

Design Fundamentals

When designing **Hbase table** we have to consider the following questions:

- What should be a **row key** and what should it contain?
- How many **column families** should a **table** have?
- What **columns** (**qualifiers**) should be included in each **column family**?
- What information should go into the **cells**?
- How many **versions** should be stored for each **cell**?

In fact **HBase table** is a four level **hierarchical structure** where a **table** consists of **rows**, **rows** consists of **column families**, **column families** consist of **columns** and **columns** consists of **versions**

If cells contain the keys then **HBase table** becomes a **network/graph structure**

Design Fundamentals

Important facts to remember:

- [Indexing](#) is performed only for a [row key](#)
- [Hbase tables](#) are sorted based on a [row key](#)
- Everything in [Hbase tables](#) is stored as [untyped sequence of bytes](#)
- [Atomicity](#) is guaranteed only at a row level and there are no multi-row transactions
- [Column families](#) must be defined at [Hbase table](#) creation time
- [Column qualifiers](#) are dynamic and can be defined at write time
- [Column qualifiers](#) are stored as sequences of bytes such that they can represent data

Design Fundamentals

Implementation of Entity type

CUSTOMER	
cnumber	ID
first-name	
last-name	
phone	
email	

```
{ "007":  
  { "CUSTOMER":  
    { "first-name": { "timestamp-1": "James" },  
      "last-name": { "timestamp-2": "Bond" },  
      "phone": { "timestamp-1": "007-007" },  
      "email": { "timestamp-1": "jb@mi6.com" }  
    }  
  }  
}
```

HBase Table

Design Fundamentals

Implementation of **one-to-one relationship**

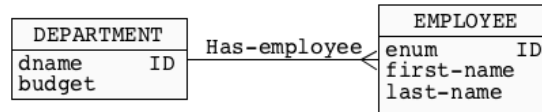


```
{ "Sales":  
  { "DEPARTMENT":  
    { "dname": { "timestamp-1": "Sales" },  
      { "budget": { "timestamp-1": "1000" }  
    }  
  }  
  { "MANAGER":  
    { "enumber": { "timestamp-2": "007" },  
      "first-name": { "timestamp-3": "James" },  
      "last-name": { "timestamp-4": "Bond" }  
    }  
  }  
}
```

HBase Table

Design Fundamentals

Implementation of **one-to-many relationship**

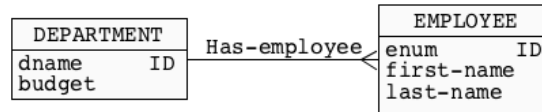


```
{ "007":  
  { "EMPLOYEE":  
    { "enumber": { "timestamp-1": "007" },  
      "first-name": { "timestamp-2": "James" },  
      "last-name": { "timestamp-3": "Bond" },  
      "department": { "timestamp-4": "Sales" }  
    }  
  }  
}
```

HBase Table

Design Fundamentals

Another implementation of [one-to-many relationship](#)

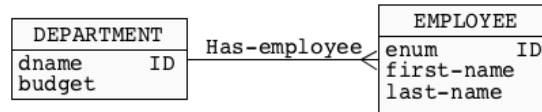


```
{ "Sales":
  { "DEPARTMENT":
    { "dname": { "timestamp-1": "Sales" },
      "budget": { "timestamp-1": "1234567" }
    }
  { "EMPLOYEE":
    { "007": { "timestamp-2": "James Bond" },
      "008": { "timestamp-3": "Harry Potter" },
      "009": { "timestamp-4": "Robin Hood" }
      ...
    }
  }
}
```

HBase Table

Design Fundamentals

Yet another implementation of **one-to-many relationship**




```
{ "Sales":
  { "DEPARTMENT":
    { "dname": { "timestamp-1": "Sales" },
      "budget": { "timestamp-1": "1000" }
    }
    { "HAS-EMPLOYEES":
      { "employees": { "timestamp-2": "007 James Bond",
                       { "timestamp-3": "008 Harry Potter",
                       { "timestamp-4": "009 Robin Hood"
                       ...
                        }
                      }
    }
  }
}
```

HBase Table

Design Fundamentals

Implementation of many-to-many relationship




```
{ "participation-1":
  { "EMPLOYEE":
    { "enumber": { "timestamp-1": "007" },
      "first-name": { "timestamp-2": "James" },
      "last-name": { "timestamp-3": "Bond" },
      "pnumber": { "timestamp-4": "project-1" },
                  { "timestamp-5": "project-2" }
      ...
    }
  }
}
```

HBase Table

Design Fundamentals

Another implementation of [many-to-many relationship](#)




```
{ "participation-1":  
  { "PROJECT":  
    { "pnumber": { "timestamp-1": "project-1" },  
      "budget": { "timestamp-2": "12345.25" },  
      "employee": { "timestamp-3": "007",  
                    { "timestamp-4": "008",  
                    { "timestamp-5": "009",  
                    ...  
                }  
            }  
  }  
}
```

HBase Table

Design Fundamentals

Another implementation of [many-to-many relationship](#)



```
{ "participation-1":  
  { "PARTICIPATION":  
    { "pnumber": { "timestamp-1": "project-1",  
                  "employee": { "timestamp-2": "employee-007" }  
    }  
  }  
}
```

HBase Table

```
{ "participation-2":  
  { "PARTICIPATION":  
    { "pnumber": { "timestamp-1": "project-1",  
                  "employee": { "timestamp-2": "employee-008" }  
    }  
  }  
}
```

HBase Table

Design Fundamentals


Note, that it is possible to group in one **Hbase table** rows of different types

```
{ "employee-007":
  { "EMPLOYEE":
    { "enumber": { "timestamp-1": "007" },
      "first-name": { "timestamp-2": "James" },
      "last-name": { "timestamp-3": "Bond" }
    }
  },
  { "project-1":
    { "PROJECT":
      { "pnumber": { "timestamp-4": "1" },
        "budget": { "timestamp-5": "12345.25" }
      }
    },
    { "participation-2":
      { "PARTICIPATION":
        { "pnumber": { "timestamp-1": "project-1" },
          "employee": { "timestamp-2": "employee-007" }
        }
      }
    }
  }
}
```

HBase Table

Design Fundamentals

Implementation of fact with dimensions




```
{ "1234567":  
  { "MEASURE":  
    { "amount": { "timestamp-1": "1000000" }  
    },  
    "BUYER":  
    { "phone": { "timestamp-1": "242214339" },  
      "first-name": { "timestamp-1": "James" },  
      "last-name": { "timestamp-1": "Bond" }  
    },  
    "SELLER":  
    { "phone": { "timestamp-1": "242215612" },  
      "first-name": { "timestamp-1": "Harry" },  
      "last-name": { "timestamp-1": "potter" }  
    }  
  }  
}
```

HBase Table

Design Fundamentals

Implementation of graph structure



```
{ "A":  
  { "R":  
    { "1": { "timestamp-1": "B" },  
      "2": { "timestamp-1": "D" }  
    },  
    "S":  
    { "1": { "timestamp-1": "C" }  
    }  
  }  
}
```


HBase Table

```
{ "B":  
  { "R":  
    { "1": { "timestamp-1": "D" }  
    }  
  }  
}
```

HBase Table

Design Fundamentals

Implementation of [graph structure](#)



```
{ "C":  
  { "S":  
    { "1": { "timestamp-1": "D" }  
    }  
  }  
}
```

HBase Table

```
{ "D":  
}
```

HBase Table

HBase Data Model

Outline

[Background](#)

[Logical view of data](#)

[Design fundamentals](#)

[Physical implementation](#)

Physical implementation

HBase is a database built on top of HDFS

HBase tables can scale up to billions of rows and millions of columns

Because Hbase tables can grow up to terabytes or even petabytes, Hbase tables are split into smaller chunks of data that are distributed across multiple servers

Chunks of data are called as regions and servers that host regions are called as region servers

Region servers are usually collocated with data nodes of HDFS

The splits of Hbase tables are usually horizontal, however, it is also possible to benefit from vertical splits separating column families

Region assignments happen when Hbase table grows in size or when a region server is malfunctioning or when a new region server is added

References

Kerzner M., Maniyam S., HBase Design Patterns, Packt Publishing 2014
(Available from UoW Library)

Jiang, Y. HBase Administration Cookbook, Packt Publishing, 2012
(Available from UoW Library)

Dimiduk N., Khurana A., HBase in Action, Mannig Publishers, 2013

Spaggiari J-M., O'Dell K., Architecting HBase Applications, O'Reilly, 2016